

Summary Table

Part	Epochs	Learning Rate	Starting model	Architecture	Train Loss	Validation Accuracy	Validation Loss	mAP
2	5	0.01, stepped	None	AlexNet	0.510573	78.9222	0.444362	n/a
4	3	0.0001, fixed	Part 2	AlexNet	0.417429	80.8514	0.417429	n/a
4	5	0.0001, fixed	Part 2	AlexNet	0.346343	81.2978	0.405009	n/a
4	8	0.0001, fixed	Part 2	AlexNet	0.485231	83.0198	0.377874	n/a
5	5	0.01, stepped	None	Fully Connected CNN	0.383046	79.7353	0.433146	n/a
5	1	1.00E-100	DetectNet	DetectNet	1.08067	n/a	0.804193	13.4135


GPU Task 1

In this task, I completed two experiments. For the first one I trained the Alexnet network for 2 epochs on 16 pictures of beagles, some “Louis” and some “not Louis”. After 2 epochs it performs at chance, classifying images of both classes as Louis with slightly above 50% certainty (Figure 1A). For the second one, I trained it for 100 epochs. This time it gets the classifications correct with at or near 100% confidence (Figure 1B). However these tests are with images from the training data rather than images the model has never seen before, and the model is likely overfitting.



Predictions

Louie	51.59%
Not Louie	48.41%



Predictions

Louie	50.67%
Not Louie	49.33%



Predictions

Louie	100.0%
Not Louie	0.0%



Predictions



Not Louie	99.62%
Louie	0.38%

B

Figure 1: A: Results after 2 epochs. Top image is of Louis, bottom image is not. B: Results after 100 epochs.

GPU Task 2

The purpose of this task was to learn about how validation data can be used to monitor a model's performance with new data while it trains. In this task I trained a neural network to detect dogs versus cats. First the data has to be preprocessed to a size compatible with the AlexNet, and organized into folders where the folder name is the category label, and with 25% of the data set aside as a validation set. Then I trained AlexNet for 5 epochs. The validation data had above 80% accuracy after the 5 epochs as shown in Figure 2, and it categorized a picture of a dog correctly with greater than 90% probability.



Predictions

dogs	92.1%
cats	7.9%

Figure 2: Dogs vs. Cats with 5 epochs. The plot shows the training and validation loss and accuracy. The image of a dog is correctly classified as a dog.

GPU Task 3

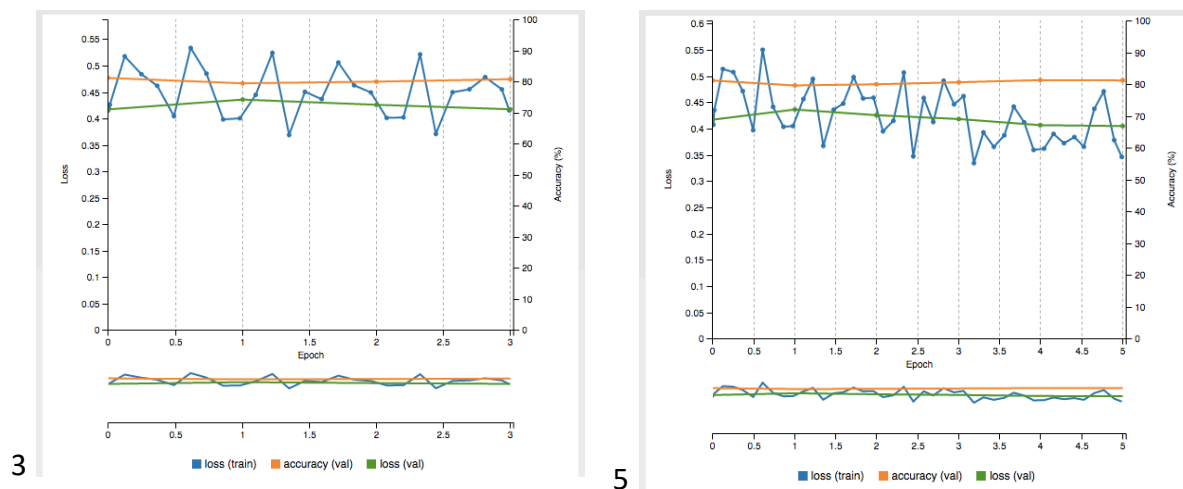
The purpose of this task was to deploy a trained model in an application. I imported the architecture and weights from the model from task 2 into a new caffe classifier object in a jupyter notebook. The notebook calculates the predicted probabilities for each class for test images, selects the argmax and prints one statement if it a cat and another if it is a dog as shown in Figure 3. This process can be consolidated into a single script that simply shows the output instead of the image with the output.



Figure 3: Example outputs for images of dogs and cats in the deployment application.

GPU Task 4

For the first part of this task, I took the model from Task 2 and trained it for more epochs to see the effects of training for a longer time. I tested 3, 5, and 8 additional epochs and found they brought only a small improvement to the validation accuracy and loss shown in the summary table. The loss and accuracy curves are in Figure 4. Then I tested using an AlexNet model pretrained on the Imagenet data set. It correctly identified a picture of Louis from Task 1 as a beagle (class 162). This can be used in an application to pair the image with its text label as shown in Figure 5.






Figure 4: Train and validation loss and accuracy curves for 3, 5, and 8 additional epochs




Figure 5: Pretrained AlexNet application output for an image of a beagle

GPU Task 5

For this task I implemented several methods of object localization. I first created an object detector that using a sliding window method to divide the image into a grid of 256x256 sections and for each section predicts whether the image is a dog or a cat. The output is a very imprecise location for the dog in the image as shown in Figure 6.

Next I took the AlexNet architecture and changed it to be a fully connected CNN. This model did a better job of detecting where in the image the dog is, but also had false positives as show in figure 7, which also includes the accuracy and loss curves.

Finally I took a DetectNet architecture and fine-tuned it for one epoch with a very slow learning rate for the dogs and cats data set. Figure 8 contains the mean average precision (mAP) and loss curves, and Figure 9 contains example outputs. This application does a much better job of

localizing the dog in the image, but in some images detects multiple overlapping dogs. This is consistent with the mAP of 13%.




Figure 6: Sliding window prediction as a method of object localization

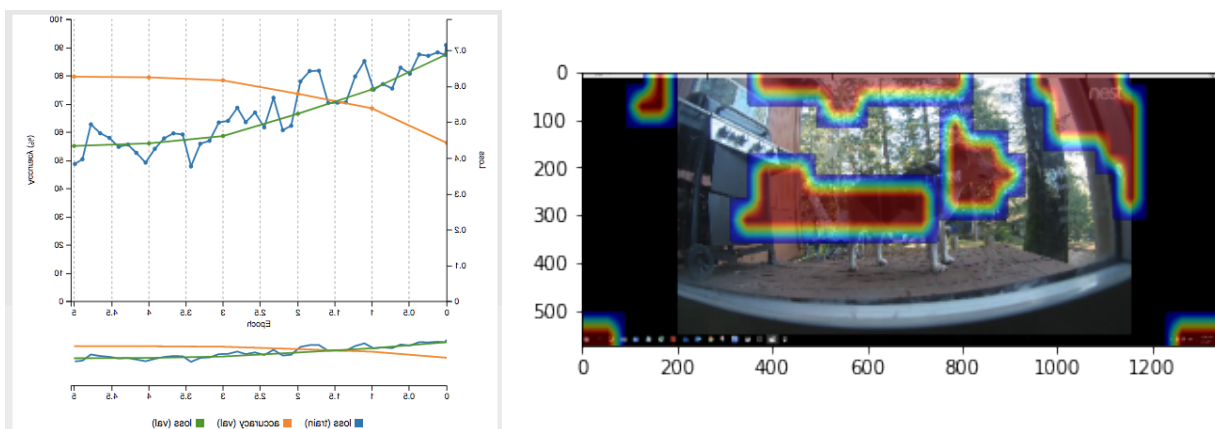


Figure 7: Fully Connected CNN Accuracy and Loss Curves and Example Output




Figure 8: DetectNet Fine-Tuning Accuracy and Loss Curves

Inference visualization



Inference visualization



Figure 9: DetectNet Example Outputs